

# SQL BOOTCAMP

*Mason Gallo*

*Instructor*

**COURSE**

---

**PRE-WORK**

---

## IF YOU HAVEN'T ALREADY...

---

- Read the README:
  - [README](#)
- Pay particular attention to the cheatsheet and exercises

**OPENING**

---

# GETTING STARTED

# SQL BOOTCAMP

## WHO AM I

- Data Scientist
- Open Source Contributor
- ML Researcher - Educational Technology
- Data Science Instructor @ GA



# SQL BOOTCAMP

---

## MY PHILOSOPHY

- If you're not sure, please ask!
- If you're still stuck, we'll revisit together
- Be considerate of your fellow students
- Participate!!!
- NO ONE knows everything
- Anything worth knowing is hard
- Breaks



# SQL BOOTCAMP

---

## LEARNING OBJECTIVES

- Learn enough SQL to effectively query data from a single table
- Learn enough SQL to effectively query data from multiple tables
- Understand where to go when you encounter new SQL problems
- Build a roadmap for your career next-steps with SQL

---

# SQL BOOTCAMP

---

## AGENDA

- Installation
- Databases 101
- Basic queries
- Aggregation
- Joins
- Dealing with data
- Next steps for your career



---

**SQL BOOTCAMP**

---

# INSTALLATION

---

# DOWNLOAD CLASS MATERIALS

---

▸ All materials are available\*\*\* [HERE](#)

Click “Clone or Download” -> “Download Zip”

Slides are available in the main directory with filename sql\_bootcamp.pdf

\*\*\*If you know git, you can always clone the repo

# INSTALLATION

---

- If you're on Windows, go [HERE](#)
- If you're on OSX, go [HERE](#)

If for some reason you're not using the recommended installation, please raise your hand!

---

# GET ACQUAINTED

---

- Take a few minutes to get acquainted with the client for your OS
- Find the area to enter “queries”
- Type into the query area:

`select * from db`

- Find the “run” or “execute” button and run the query
- What’s the output?

**SQL BOOTCAMP**

---

# DATABASES 101

---

# **YOUR THOUGHTS**

---

**How do you store your data? At work? Personal?**

---

# MY THOUGHTS

---

**Storing social performance for many brands over several years**

---

# WHAT ARE THE OPTIONS

---

- Storing in a raw text file
- Storing in excel or google sheets



---

# **LIMITATIONS**

---

- Slow and inefficient
- Hard to share with coworkers
- Difficult to use with other applications
- Doesn't scale (millions? billions?)

---

# WHAT DO WE WANT

---

- A methodology that can be used by many people and machines
- Scalability
- Speed
- Safety

---

# WHAT DO WE WANT

---

Databases are a **structured** data source optimized for efficient **retrieval** and **storage**

---

# WHAT DO WE WANT

---

Databases are a **structured** data source optimized for efficient **retrieval** and **storage**

**Structured:** we must pre-define the structure of our data

---

# WHAT DO WE WANT

---

Databases are a **structured** data source optimized for efficient **retrieval** and **storage**

**Structured:** we must pre-define the structure of our data

**Retrieval:** we must have a systematic way to retrieve the structured data

---

# WHAT DO WE WANT

---

Databases are a **structured** data source optimized for efficient **retrieval** and **storage**

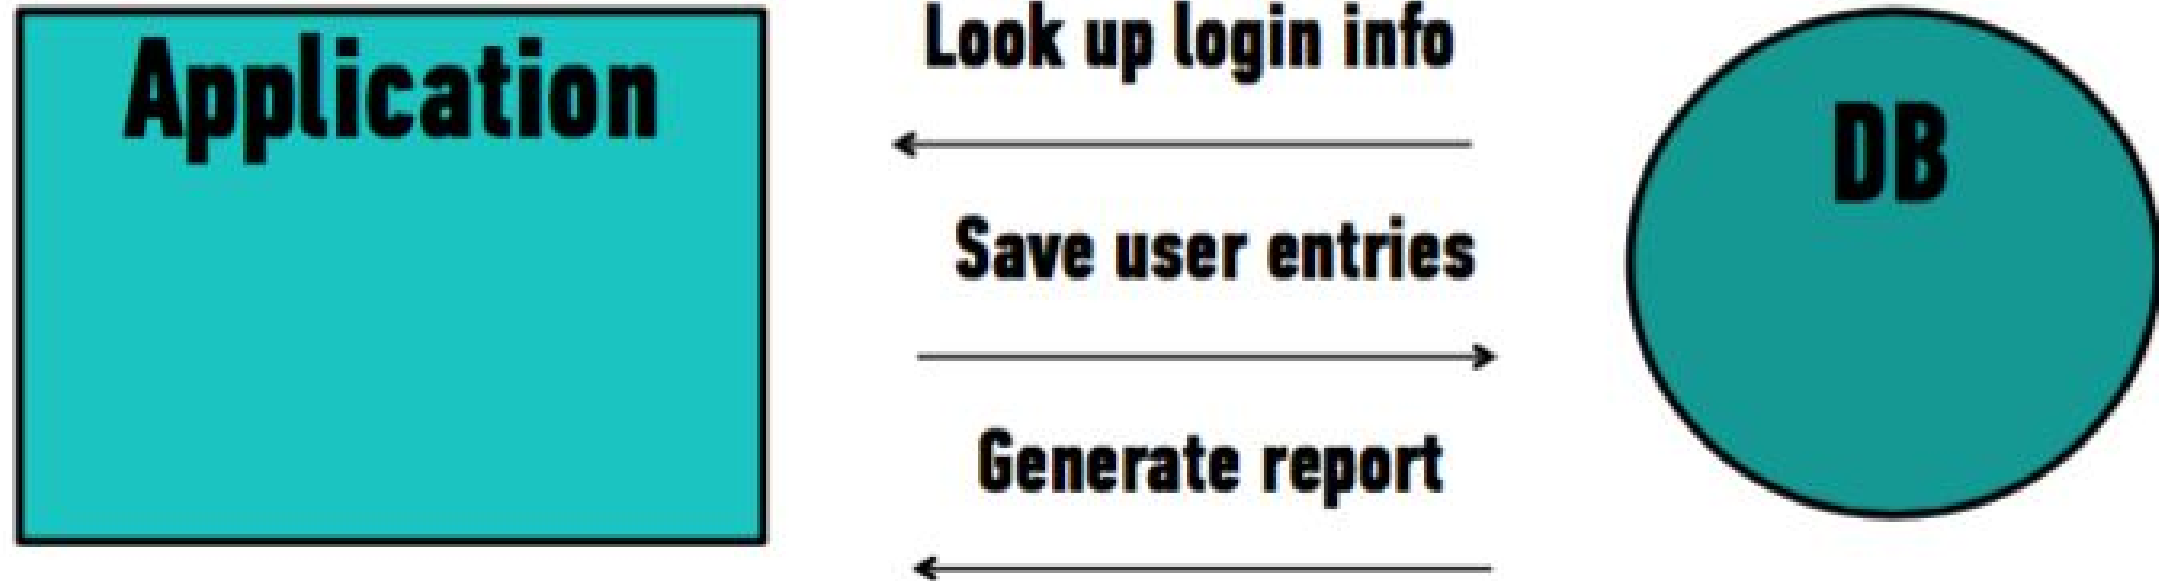
**Structured:** we must pre-define the structure of our data

**Retrieval:** we must have a systematic way to retrieve the structured data

**Storage:** we must have the ability to read and save data

# YEAH YEAH YEAH, SO WHERE DOES SQL FIT IN?

**SQL is the language used to perform these actions!**



---

# WHAT IS SQL?

---

SQL (Structured Query Language) is a query language designed to **Extract, Transform, Load** data in relational databases



---

# WHO USES SQL?

---

Web Developers

Data Analysts

Data Scientists

Product Managers

Statisticians

System Administrators

Backend Developers

...And many more

---

# THE TERMS NO ONE EXPLAINS

---

- Database: think of as a directory or folder of tables
- Table: a collection of related structured data, usually as rows / columns
- Record: SQL lingo for “row”
- Field: SQL lingo for “column”
- Localhost: this computer (ie the computer you’re using)
- Root: a user with FULL permission to do virtually anything
- UTF-8: capable of encoding all unicode characters

**SQL BOOTCAMP**

---

# BASIC QUERIES

---

## NOTE BEFORE WE GET STARTED

---

- Capitalization doesn't matter (SQL isn't case-sensitive)
- SQL ignores whitespace (but don't go crazy)
- Be OCD about your punctuation (it matters!)

---

# CREATE YOUR FIRST DATABASE

---

```
CREATE DATABASE ga;
```

```
USE ga;
```

Now, use your GUI to import users.sql, selecting UTF-8 if asked

```
mysql -u root -p test ga < users.sql
```

---

# BASIC QUERIES

---

SELECT \*  
FROM users

*Return all rows from all  
columns*

---

# BASIC QUERIES

---

```
SELECT title, first_name,  
last_name  
FROM users
```

*Return all rows from specific  
columns*

# BASIC QUERIES

---

```
SELECT DISTINCT title,  
first_name, last_name  
FROM users
```

*Return unique rows from the  
specified columns*



# BASIC QUERIES

---

SELECT \*  
FROM users  
WHERE <condition>

*Return all columns under a  
specified condition*

# BASIC QUERIES

---

SELECT \*

FROM users

WHERE state = “arizona”

# BASIC QUERIES

---

```
SELECT *  
FROM users  
WHERE state LIKE "arizona"  
AND title LIKE "Miss"
```

---

## \*\*\*BONUS\*\*\*

---

```
SELECT *  
FROM users  
WHERE state LIKE "arizona"  
AND title LIKE binary "miss"
```

Binary will force case  
sensitivity!

---

# BASIC QUERIES

---

```
SELECT *  
FROM users  
WHERE state IN  
("arizona", "florida")
```

# BASIC QUERIES

---

```
SELECT *  
FROM users  
WHERE state IN  
("arizona", "florida")
```

What happens if we put a  
NOT in front of IN?

# BASIC QUERIES

---

```
SELECT *  
FROM users  
WHERE zip = 10007
```

Notice the use of = here  
instead of LIKE. Any ideas?

We can also use >, <, !=, >=,  
<=

# BASIC QUERIES

---

SELECT \*

FROM users

WHERE street LIKE “%rd”

% matches anything

\_ matches a single character



# BASIC QUERIES

---

```
SELECT *  
FROM users  
WHERE street LIKE "%rd"  
ORDER BY first_name
```

What happens if you put  
DESC after first\_name?

# BASIC QUERIES

---

```
SELECT *  
FROM users  
WHERE street LIKE "%rd"  
ORDER BY first_name  
LIMIT 5
```

What does LIMIT do?

---

# YOUR TURN

---

Complete the warm up exercises [here](#)

**SQL BOOTCAMP**

---

# AGGREGATION

---

## YOUR TURN

---

What if we want to compute overall statistics about the data?

Think: we want counts, averages, etc

---

# AGGREGATION

---

SELECT \*  
FROM users

*Return all rows from  
columns*

---

# AGGREGATION

---

```
SELECT count(*)  
FROM users
```

*Return the count of all rows  
from all columns*

---

# AGGREGATION

---

```
SELECT distinct first_name  
FROM users
```

*Return the unique first  
names*



---

# AGGREGATION

---

```
SELECT count(distinct first_name)  
FROM users
```

*Return the count of unique  
first names*

# AGGREGATION

---

```
SELECT avg(zip)  
FROM users
```

*Return the average zip code  
(yes, I know this doesn't  
make sense)*

---

# AGGREGATION

---

```
SELECT avg(zip) as “nonsense”  
FROM users
```

*Rename our avg zip to  
“nonsense”*

# AGGREGATION

---

SELECT gender, avg(zip) as  
“nonsense”

FROM users

GROUP BY gender

*What if we want to know the  
avg zip by gender?*

*We can use sum, avg, min,  
max, count*

# AGGREGATION

---

SELECT state, gender, avg(zip) as  
“nonsense”

FROM users

GROUP BY state, gender

*What happened here?*

# AGGREGATION

---

SELECT state, gender, avg(zip) as  
“nonsense”

FROM users

GROUP BY state, gender

ORDER BY 1

*What happened here?*

# AGGREGATION

---

SELECT state, gender, avg(zip) as  
“nonsense”

FROM users

GROUP BY state, gender

ORDER BY 1

*What happened here?  
Shorthand?*

# AGGREGATION

---

SELECT state, gender, avg(zip) as  
nonsense  
FROM users  
GROUP BY state, gender  
HAVING nonsense > 50000

*What happened here? What  
happens if we use WHERE  
instead?*



---

# AGGREGATION

---

**SELECT** <columns>  
**FROM** <table>  
**WHERE** <condition>  
**GROUP BY** <columns>  
**HAVING** <condition on aggregates>  
**ORDER BY** <columns>  
**LIMIT** <number>

*General SQL structure  
(putting it all together)*

---

# YOUR TURN

---

Complete the aggregation exercises [here](#)

---

**SQL BOOTCAMP**

---

# JOINS

# JOINS

---

What if we want to deal with multiple tables?

Think: we want to capture information from 2 or more tables simultaneously

Import flights.sql

# JOINS

---

A **relational database** is organized in the following manner:

- ▶ A database has **tables** which represent individual entities or objects
- ▶ Tables have a predefined **schema** – rules that tell it what columns exist and what they look like

# JOINS

---

A **relational database** is organized in the following manner:

► **table**

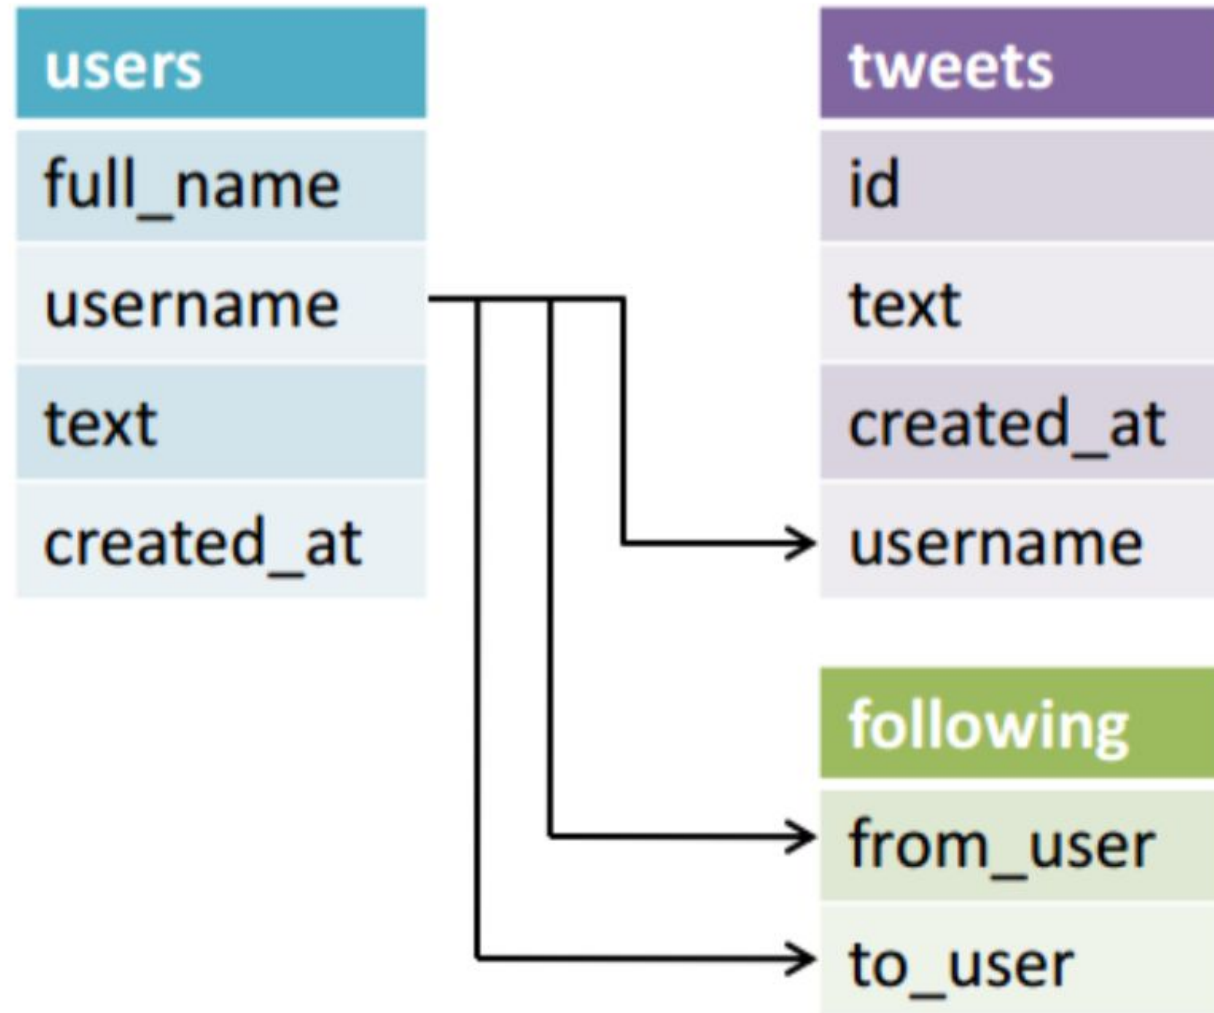
<b>id</b>	<b>first name</b>	<b>last name</b>	<b>date of birth</b>
312	Joe	Smith	1980-12-24
1532	Michelle	Anderson	1973-03-12

► **schema**

id	bigint
first_name	char(36)
last_name	char(36)
date_of_birth	timestamp


# JOINS

---



# JOINS

---





# JOINS

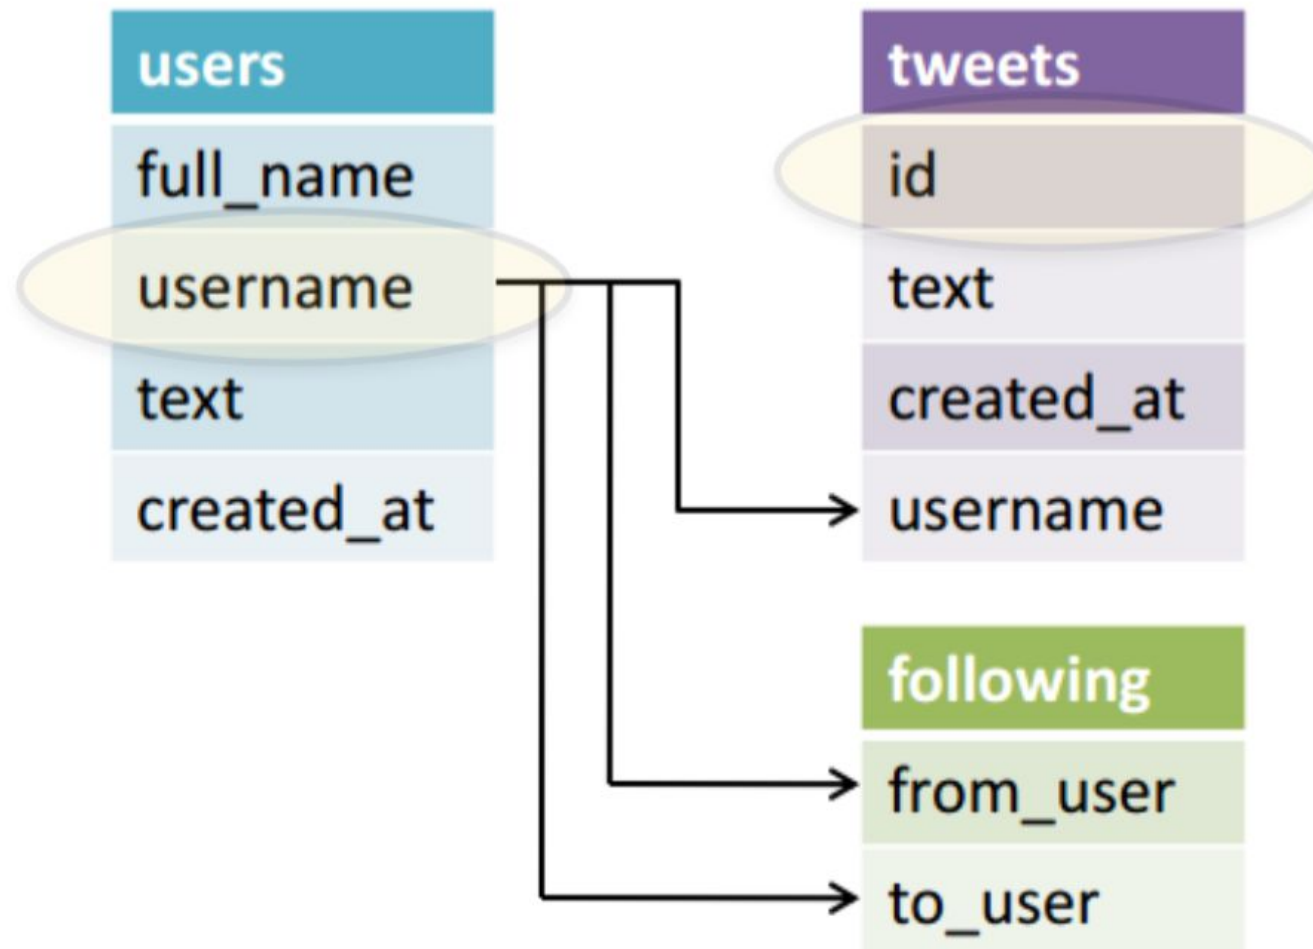
---

Each table should have a PRIMARY KEY

Think: a unique identifier for each row

Ex: SSN

# JOINS



---

# JOINS

---

Each table should have a PRIMARY KEY


Think: a unique identifier for each row

Ex: SSN

Additionally, each table may have a FOREIGN KEY

Think: an ID that links one table to another

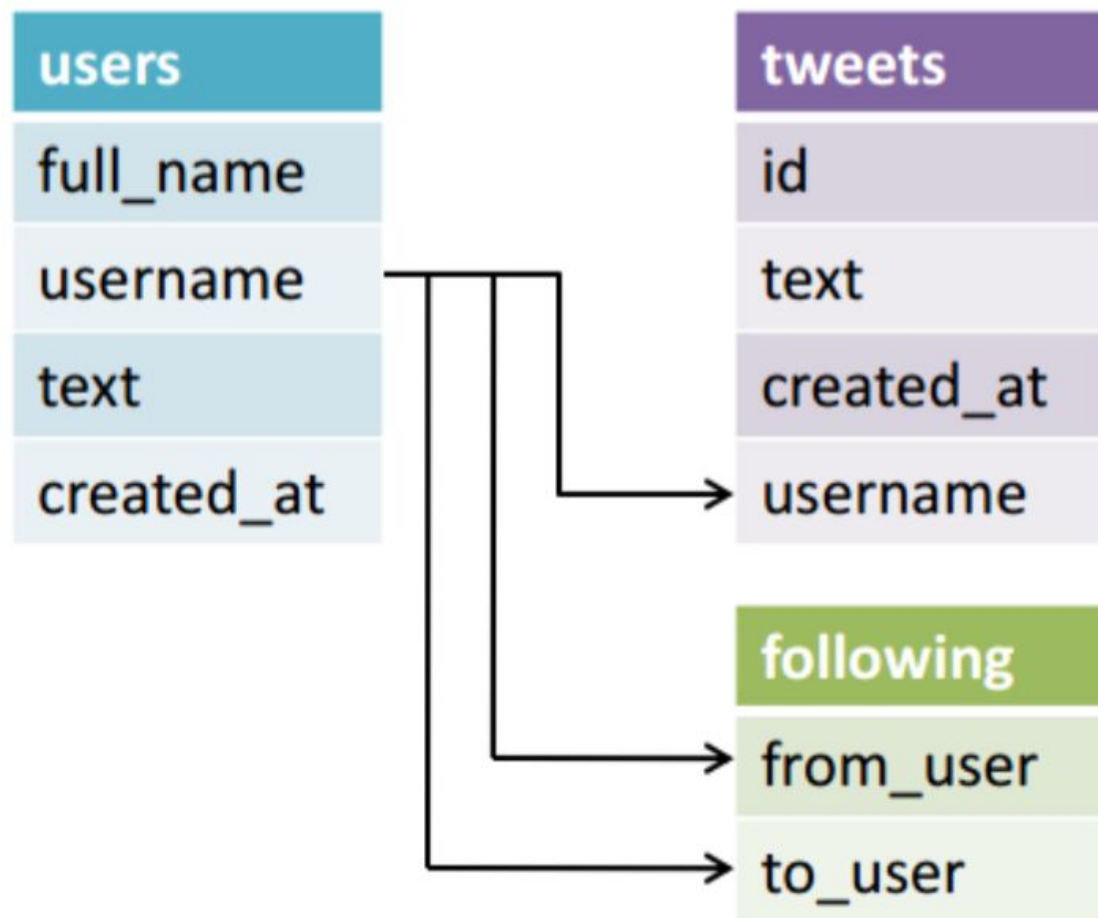
# JOINS



# JOINS

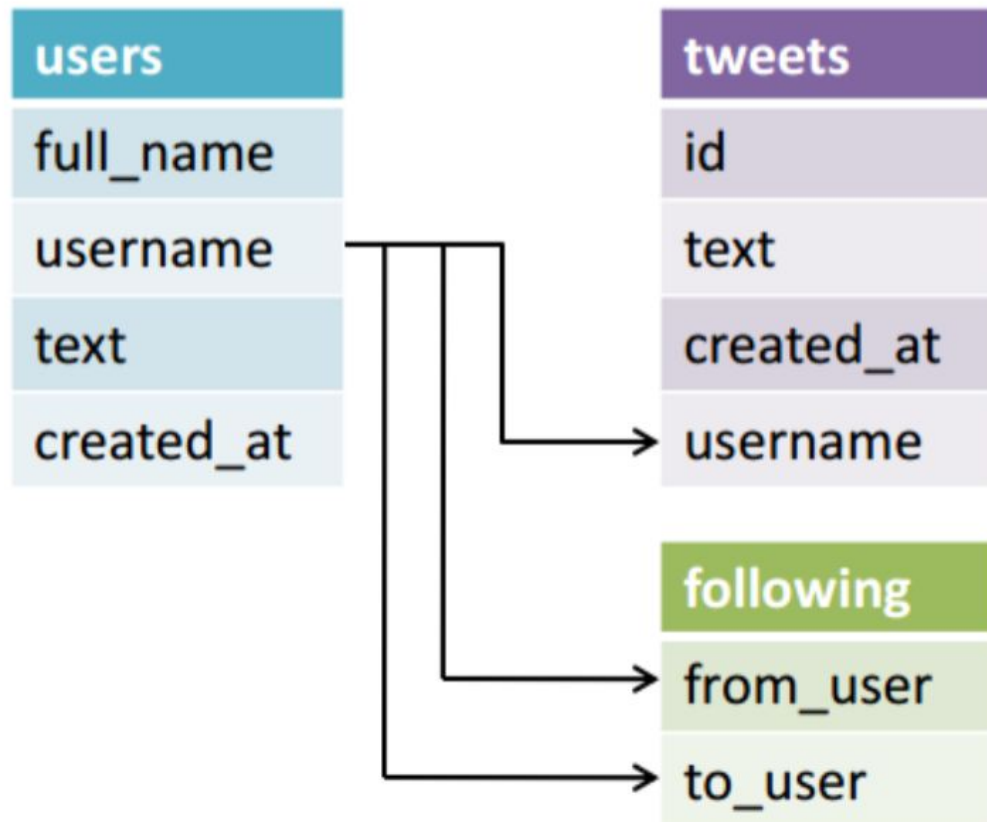
---

*Create a table with all the users' full names and their tweets*



# JOINS

*Create a table with all the users' full names and their tweets*

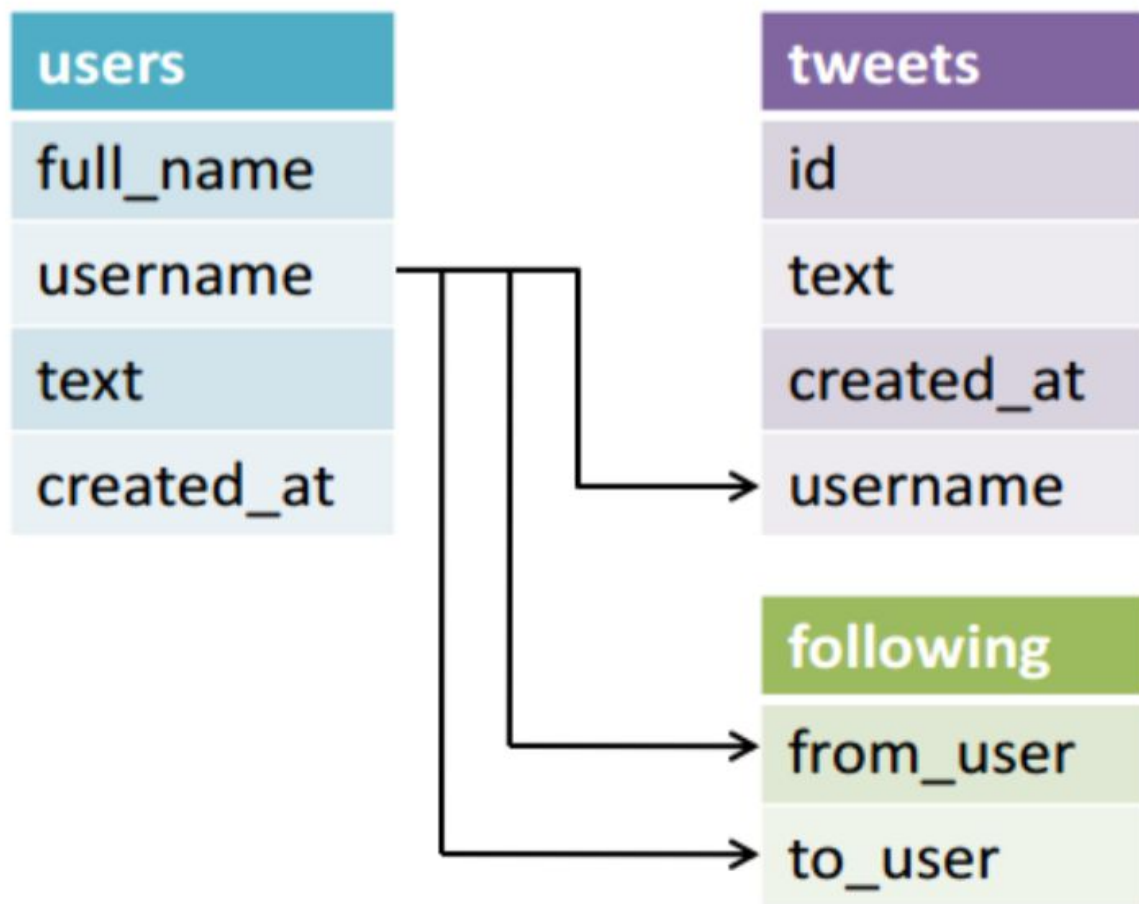


<u>full_name</u>	<u>tweet</u>
Joe Smith	Hello, world!
Joe Smith	Just tweetin'
Michelle Ng	I am eating pizza tonight

# JOINS

---

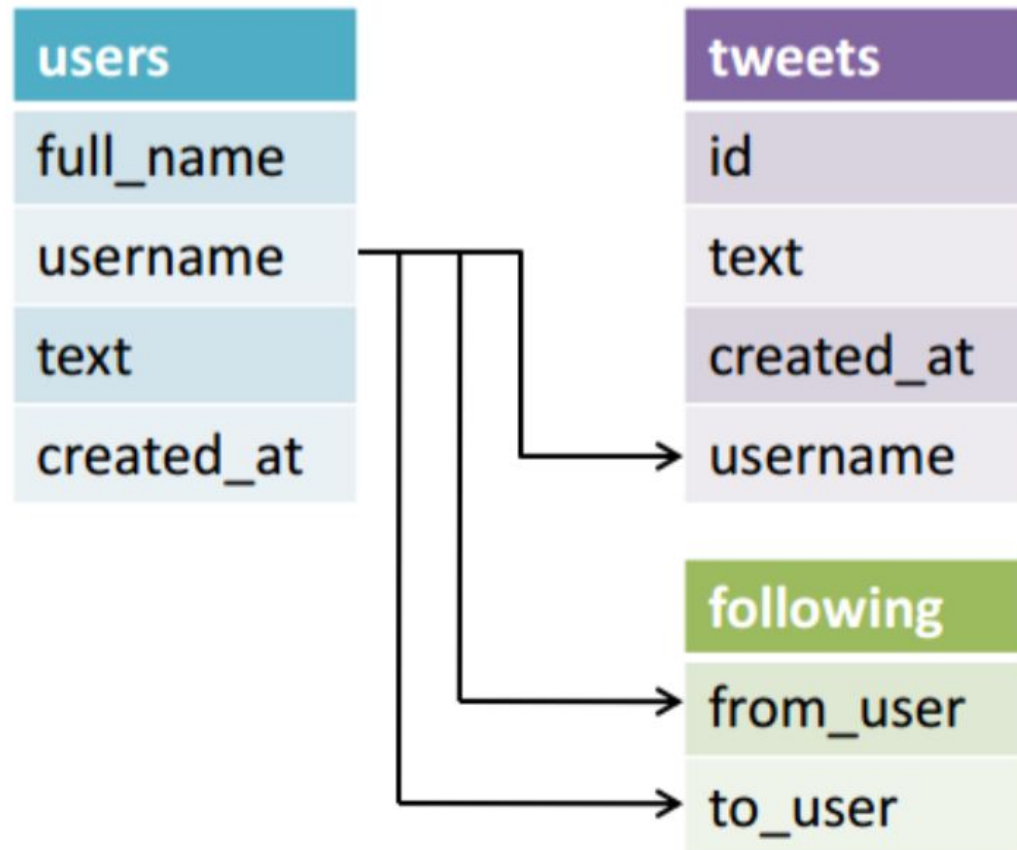
*Create a table with all the users' full names and their tweets*



```
SELECT
  users.full_name,
  tweets.text
```

# JOINS

*Create a table with all the users' full names and their tweets*

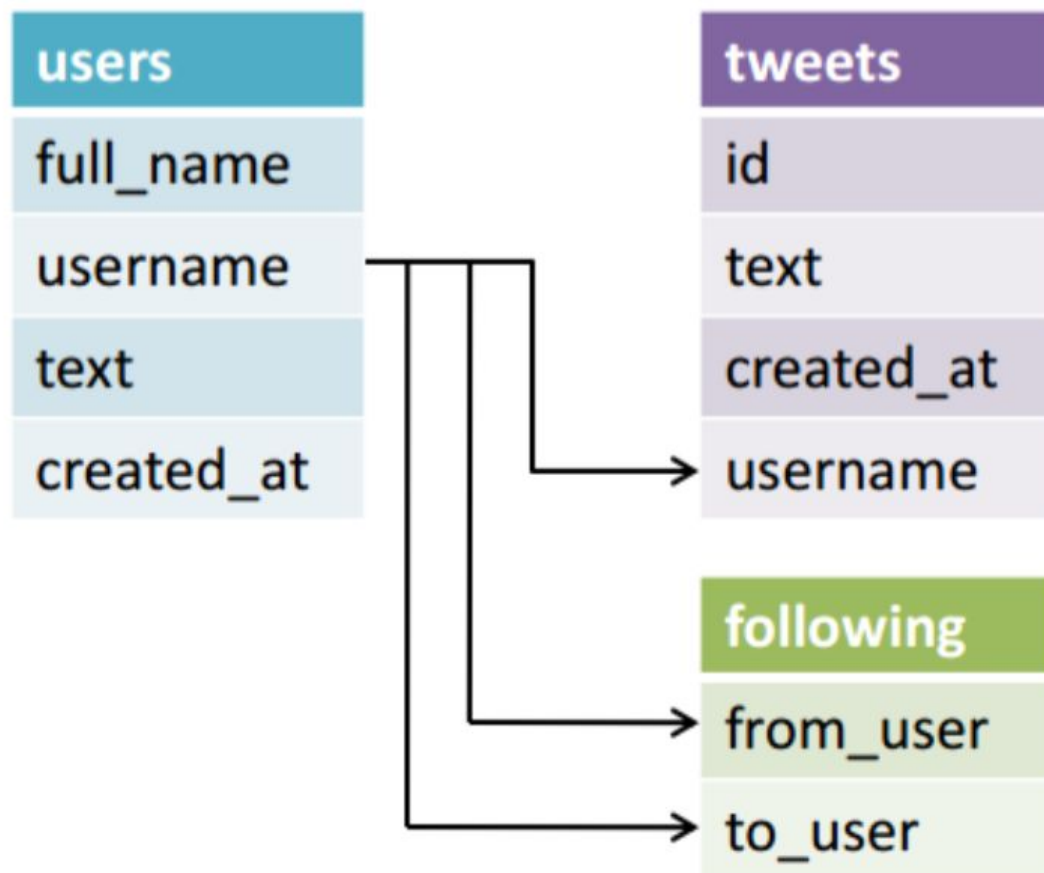


```
SELECT
  users.full_name,
  tweets.text
FROM users
JOIN tweets
ON users.username = tweets.username
```



# JOINS

*Create a table with all the users' full names and their tweets*



*Will users who never tweeted appear in the list?*

# JOINS

---

**JOIN** *will only include entries that occur in both tables.*

```
SELECT
  full_name,
  text
FROM users
JOIN tweets
ON users.username = tweets.username
```

<u>full_name</u>	<u>tweet</u>
Joe Smith	Hello, world!
Joe Smith	Just tweetin'
Michelle Ng	I am eating pizza tonight

# JOINS

---

**LEFT JOIN** *will always include all entries from the left table, even if there are no matches in the other table.*

```
SELECT
  full_name,
  text
FROM users
LEFT JOIN tweets
ON users.username = tweets.username
```

<u>full_name</u>	<u>tweet</u>
Joe Smith	Hello, world!
Joe Smith	Just tweetin'
Michelle Ng	I am eating pizza tonight
Jim Rogers	

# JOINS

---

**FULL OUTER JOIN** *will always include all entries from both tables, even if there are no matches in the other table.*

```
SELECT
  full_name,
  text
FROM users
FULL OUTER JOIN tweets
ON users.username = tweets.username
```

<u>full_name</u>	<u>tweet</u>
Joe Smith	Hello, world!
Joe Smith	Just tweetin'
Michelle Ng	I am eating pizza tonight
Jim Rogers	OK, deleting my account



# JOINS

---

The holes in the resulting table are called **NULLs**.

**NULL** indicates missing data.

Note that **NULL** is not the same as zero or an empty string "", it really means that there is no data.

<u>full_name</u>	<u>tweet</u>
Joe Smith	Hello, world!
Joe Smith	Just tweetin'
Michelle Ng	I am eating pizza tonight
Jim Rogers	<b>NULL</b>
<b>NULL</b>	OK, deleting my account

# JOINS

---

For example, to print a list of users without tweets, we'd write


```
SELECT full_name  
FROM users  
FULL OUTER JOIN tweets  
ON users.username = tweets.username  
WHERE tweets.text IS NULL
```




<u>full_name</u>
Jim Rogers

# JOINS


## SQL JOINS




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```



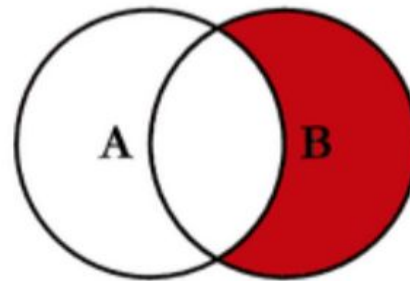
```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL  
OR B.Key IS NULL
```

# JOINS

---

```
SELECT <columns>  
FROM <table>  
JOIN <otherTable>  
ON <table.key> = <otherTable.key>  
JOIN <yetAnotherTable>  
ON <otherTable.key> = <yetAnotherTable.key>
```

*General SQL structure*

## NOTE

You can combine as many **JOINS** as you want!



---

# JOINS

---

*General SQL structure*

**SELECT <columns>**  
**FROM <table>**  
**[INNER|LEFT|RIGHT|FULL OUTER] JOIN <otherTable>**  
**ON <table.key> = <otherTable.key>**  
**WHERE <condition>**  
**GROUP BY <columns>**  
**HAVING <condition>**  
**ORDER BY <columns> [DESC|ASC]**  
**LIMIT <number>**

**SQL BOOTCAMP**

---

# DEALING WITH DATA

---

# DEALING WITH DATA

---

SELECT 1+1

What just happened?

# DEALING WITH DATA

---

SELECT 1+1

SELECT “hi”

What just happened?

---

# DEALING WITH DATA

---

```
SELECT first_name, last_name,  
CASE WHEN last_name like "n%" THEN "Starts with n" ELSE "Doesn't"  
END  
FROM users
```

Think: if...then statement!

---

# DEALING WITH DATA

---

```
SELECT first_name, last_name,  
FROM users  
WHERE first_name in (  
    SELECT DISTINCT first_name  
    FROM presidents  
)
```

This is called a subquery!

---

# DEALING WITH DATA

---

Check out the cheatsheet [HERE](#)

**SQL BOOTCAMP**

---

**NEXT STEPS**



---

# NEXT STEPS FOR WEB DEVELOPERS

---

Install the [LAMP Stack](#)

You now know the M in LAMP!

Try creating a site and using SQL to store your data

---

# NEXT STEPS FOR DATA ANALYSTS / DATA SCIENTISTS

---

Get a toy dataset

Use in conjunction with a data visualization tool like Excel or Tableau

Excel has [Power Query](#) and Tableau has native SQL support!

---

## NEXT STEPS FOR BACKEND DEVELOPERS

---

Get a toy dataset and manipulate using your favorite scripting language, like Python (I'm biased!)

---

# OPINION

---

Be weary of people telling you that SQL is “too slow” or “too fast”

Make sure they’re not trying to sell you something!

# OPINION

---

You should not be “cleaning” your data in SQL

You should not be “formatting” your data in SQL for visualization

SQL is only for storing and retrieving raw data!

---

# OPINION

---

You now know probably 80% of the SQL you need

Many people will never need to use the remaining 20% EVER

Make sure your skills don't atrophy!

# KEEP IN TOUCH

---

If you enjoyed my class, I am starting a blog:

Follow me on [Twitter](#)

If you're on Github:

Follow [Me](#)

If you're on LinkedIn:

Connect with [Me](#)